

# Iterative Linear Quadratic Regulator in C++

Thomas Lowe

March 9, 2016

If you want to move a robot arm or control a vehicle or any machine then a great way to do it is with feedback control. You simply specify the dynamics of your machine with the environment:

$$\dot{x} = f(x, u)$$

where  $x$  are your degrees of freedom, such as robot arm joint angles and angular velocities.  $u$  are your control signals, such as joint torques, and  $f$  is the nonlinear dynamics. You then define the running cost of your state and signals:  $l(x, u)$  and the final cost for the final state of your machine's trajectory  $l_f(x)$ :

$$\sum_{k=0}^K l(x_k, u_k) + l_f(x_K)$$

We then want to automatically find an optimal controller that minimises the cost subject to the dynamics. An Iterative Linear Quadratic Regulator (ILQR) is a useful means of solving this Optimal Feedback Control problem. It generates a feedforward (open loop) signal  $u_0$  and a feedback (closed loop) term:

$$u = u_0 + L(x_0 - x)$$

One could of course have just found an optimal open loop signal  $u_0$  but the feedback term allows perturbations from the assumed dynamics  $f$  to be dealt with, and also provides a higher update-rate controller in the case that ILQR is applied intermittently, e.g. in a receding-horizon Model Predictive Control setting. So Optimal Feedback Control is more robust than just Optimal Control.

There are few libraries that do ILQR and particularly not in C++. So I have [implemented the algorithm in C++](#) based off the theory and algorithms by Mitrovic 2010 [1], Tassal IROS 2012 [2], Tassal ICRA 2014 [3] and StudyWolf 2016 [4]. The implementations are very similar but where there are differences I have used macros to allow the user to try the different options. The default is what gave best results for my tests. The parameter names also vary, I have referred to the names that are used in the different papers but I have chosen to use the syntax of standard [LQRs as used in wikipedia](#), this is clearest to me.

I have kept the library very lightweight, it only relies on the [Eigen vector library](#) and is implemented just in the header file ILQR.h. You simply derive this ILQR class and override the system dynamics function  $f(x, u)$  and the cost functions  $l(x, u)$ ,  $l_f(x)$ . Internally these are converted into partial derivatives using finite differences, however if you know the partial derivatives of any of these functions analytically then you can override these `getDerivatives()` functions directly and save CPU cost.

Now we can test the ILQR, I have included `acrobot.h` which refers to a toy problem described in 1991 by Richard Murray [5], how do you get a double pendulum to swing up and balance vertically when you only control the mid joint? This is a bit like an acrobat on the parallel bars swinging up to a hand-stand. To solve we only need to penalise difference from the desired end state and a squared torque running cost to avoid torque spikes. Try adjusting these parameters in  $l()$  and  $l_f()$  to see how they effect the results. To keep the library simple I don't animate the acrobot, just print out the trajectory, which you can load in any spreadsheet:




Figure 1: red: actuated joint, blue: unactuated joint

In addition I also include a continuous time version of the solver, which instead interprets the running cost as an integral:

$$\int_{k=0}^K l(x_k, u_k) dk + l_f(x_K)$$

This is mainly useful theoretically as it makes a simpler formulation.

## References

- [1] [Stochastic Optimal Control with Learned Dynamics Models](#) Appendix A

- [2] [Synthesis and Stabilization of Complex Behaviors through Online Trajectory Optimization](#)
- [3] [Control-Limited Differential Dynamic Programming](#)
  - [implementation](#)
- [4] [the iterative linear quadratic regulator method](#)
  - [implementation](#)
- [5] [A Case Study in Approximate Linearization: The Acrobot Example](#)
  - [equations](#)